

Deep Generative Modelling

Project Presentation for CS772A

Archit Sharma (14129)

Abhinav Agrawal (14011)

Anubhav Shrivastava (14114)

Table of contents

1. Introduction
2. Background
3. Approach
4. Results
5. Ongoing Work

Introduction

- Classic task of estimating the density function.
- Objectives in context of Deep Learning context slightly relaxed: Generate realistic samples and provide likelihood measurements.
- Advances in Deep Learning have provided a real jump in our capacity to model and generate from multi-modal high-dimensional distributions, particularly those associated with natural images.

- Represents our ability to manipulate high dimensional spaces, and extract meaningful representations
- Extremely important in context of Semi-supervised Learning (in general when training data available is low) or when we have missing data
- Naturally handles Multi-modal outputs

Some Recent Approaches

Two game changing works:

- **Variational Autoencoders:** Explicit density measurement with approximate posterior maximization.
- **Generative Adversarial Networks:** Implicit Density Maximization.

Some other frameworks based on Maximum Likelihood Estimation:
Real NVP, PixelRNN.

We looked at many frameworks: Generative Latent Optimization (GLO), Bayesian GANs, Normalizing and Inverse Autoregressive Flows.

Background

Normalizing Flows

Traditionally, variational inference employs simple families of posterior approximations to allow efficient inference. With the help of normalizing flows, a simple initial density is transformed into a density of desired complexity by applying a sequence of transformations.

- Suppose z has a distribution $q(z)$ and $z' = f(z)$, then distribution of z' is given by:

$$q(z') = q(z) \left| \det \left(\frac{\partial f^{-1}}{\partial z'} \right) \right| = q(z) \left| \det \left(\frac{\partial f}{\partial z} \right) \right|^{-1}$$

- Above mentioned simple maps can be combined several times to construct complex densities:

$$z_K = f_K \circ \dots \circ f_2 \circ f_1(z_0)$$

$$\ln q_K(z_K) = \ln q_0(z_0) - \sum_{k=1}^K \ln \left| \det \frac{\partial f_k}{\partial z_{k-1}} \right|$$

Normalizing Flows

- Normalizing flows proposes to use the following transformation:

$$f(z) = z + \mathbf{u}h(\mathbf{w}^T x + b)$$

- The determinant of the Jacobian:

$$\psi(z) = h'(\mathbf{w}^T z + b)\mathbf{w}$$

$$\left| \det \frac{\partial f}{\partial z} \right| = |\det(\mathbf{I} + \mathbf{u}^T \psi(z)^T)| = |1 + \mathbf{u}^T \psi(z)|$$

- We can apply a sequence of above transformations to get q_K :

$$\ln q_K(z_K) = \ln q_0(z_K) - \sum_{k=1}^K \ln |1 + \mathbf{u}_k^T \psi_k(z_{k-1})|$$

Real NVP transformations is a framework for doing invertible and efficiently learnable transformations, leading to an unsupervised learning algorithm with exact log-likelihoods, efficient sampling and inference of latent variables.

- Change of variable formula:

$$p_X(x) = p_Z(f(x)) \left| \det \left(\frac{\partial f(x)}{\partial x} \right) \right|$$

- Coupling layers:

$$y_{1:d} = x_{1:d}$$

$$y_{d+1:D} = x_{d+1:D} \odot \exp(s(x_{1:d})) + t(x_{1:d})$$

- Jacobian of the above transformation is a lower triangular matrix which reduces the computation cost for its calculation.

- The above transformation is invertible:

$$x_{1:d} = y_{1:d}$$

$$x_{d+1:D} = \left(y_{d+1:D} - t(y_{1:d}) \right) \odot \exp\left(-s(y_{1:d}) \right)$$

- The above mentioned transformations leaves some of the components unchanged. The **coupling layers** can be composed in alternate fasion to solve this issue.

Approach

Normalizing Flows provides a framework, incorporated within VAEs, where more complex posteriors can be obtained by using invertible transformations. The constraint on the transformation: *Determinant of Jacobian matrix should be efficiently computable*. We propose to

use Real NVP transformations. These transformations are much more powerful than those proposed in Normalizing flows, but at the same time have efficient Jacobian computation as well.

Framework Details

We look to model the Binarized MNIST. The model structure is similar to those in VAEs and Normalizing Flows.

- *Encoder*: Passes images through a set of convolutional and pooling layers. Then uses a few fully connected layers to convert each image into a fixed size embedding.
- *Transformations*: In lines with Normalizing Flows, the embedding from the encoder is passed through a sequence of real NVP transformations
- *Decoder*: The transformed embedding is converted into an image by passing through a sequence of transposed convolutional layers.

Each transformation is a a set of "coupling layers", such that no dimension of the embedding is untransformed


$$\begin{aligned}F(x) &= \mathbb{E}_{q_\phi(z|x)}[\log q_\phi(z|x) - \log p(x, z)] \\&= \mathbb{E}_{q_0(z_0)}[\log q_K(z_K) - \log p(x, z_K)] \\F(x)_{NF} &= \mathbb{E}_{q_0(z_0)}[\log q_0(z_0)] - \mathbb{E}_{q_0(z_0)}[\log p(x, z_K)] \\&\quad - \mathbb{E}_{q_0(z_0)} \left[\sum_{k=1}^K \ln |1 + \mathbf{u}_k^T \psi_k(z_{k-1})| \right] \\F(x)_{rNVP} &= \mathbb{E}_{q_0(z_0)}[\log q_0(z_0)] - \mathbb{E}_{q_0(z_0)}[\log p(x, z_K)] \\&\quad - \mathbb{E}_{q_0(z_0)} \left[\sum_{k=1}^K (s_{1,k}(b \odot z_{k-1})) \right]\end{aligned}$$

Results

Table 1: NF: Normalizing Flows, **rNVP:** Real NVP. k denotes the number of transformations

Models	$\log p(x z)$
rNVP ($k = 2$)	60.57
rNVP ($k = 5$)	75.56
rNVP ($k = 10$)	75.01
rNVP ($k = 20$)	81.37
NF ($k = 4$)	65.5
NF ($k = 10$)	68.9
NF ($k = 20$)	75.4
NF ($k = 40$)	83.4

Latent Space Interpolation



(a) Latent Space Interpolations for a simple VAE

(b) Latent Space interpolations with rNVP transformations

Figure 1

Ongoing Work

- Implement the above algorithms to larger datasets (SVHN, CIFAR10, CelebA)
- Include better stabilization techniques such as weight normalization and batch normalization in realNVP transformations to train deeper networks
- Experiment with Convolutional Neural Networks for transformations (s and t can both be any arbitrary transformation)

Questions?